Note: Forthcoming in *CLCWeb: Comparative Literature and Culture*.
Not for circulation or citation without permission of the author.

Who Speaks as AI?

Abstract: This essay puts into conversation disparate fields of language theory to address the question, "Who speaks as AI?" First, it examines the structure and composition of Large Language Models (LLMs) alongside concepts of human language structure as explored by Ferdinand de Saussure and Jacques Derrida. Second, it investigates whether LLM output possesses communicative intent, drawing on the works of Emily Bender and her co-authors as well as the "wave poem" debate initiated by Steven Knapp and Walter Benn Michaels in *Critical Inquiry*. Finally, it proposes a "paranoid" hermeneutic for interpreting AI output, after Eve Sedgwick's critique of paranoid reading, in order to expose the diffuse corporate intentionality and capitalist relations embedded into AI language.

Who says "I" when AI speaks? Almost every word in the question is under contention. The term AI is marketing, says Emily Bender, computational linguist; it implies an intelligence these large language models (LLMs)[1] do not have ("Opening Remarks"). Whether AI chatbots can speak at all, or engage in *parole*, which implies a promise, is also undetermined. But the majority of us do speak with, or to, or at, or at least *about* AI, now on an ever more frequent basis, and it's worth considering what we can know about our interlocuter and their language. Do they speak the same language as us? Do they speak for themselves, and if not, on whose behalf? Do they speak with intention, and if so, whose? These questions become more pressing when considering recent trends, like AI users becoming delusional after speaking with AI and attributing to it an other-worldly or omniscient intelligence (Hill) or isolating themselves in the midst of mental health crises, believing that AIs offers adequate support for their problems (Reiley).

The animating question for this essay opens up long-debated subjects in language theory, such as what constitutes meaning, intentionality, communication, and even language itself. But it also touches on questions of predatory capitalism, power relations, and marketing. I argue that these two threads are more intertwined than has thus far been theorized. Throughout this discussion, I put into conversation siloed areas of language theory, including computational linguistics, sociolinguistics, ordinary language philosophy, structuralism, post-structuralism, and anti-post-structuralism. Many of these are now coming together in the new field of Critical AI studies, as compiled in the Duke University Press journal *Critical AI* and elsewhere. But few of the articles I cite fundamentally take up just how implicated corporate owners are in the output of AI. I argue for a hermeneutic attention to this aspect of AI structure and affect, offering a

[1] Note: In this essay, I'll use the terminology AI to refer to the popular understanding of AI characters like ChatGPT or Claude. When referring to the technicalities of LLM structure, I will use the term LLM. Ultimately, they refer to the same thing, given that I am only discussing language-based AIs.

methodology for what I call "paranoid" readings of AI output, after Eve Sedgwick's critique of paranoid reading. I take up this task in three parts: (1) "Input," explores the question of whether we are speaking the same language as AI by comparing AI structure with theories of language; (2) "Output and Interpretation," theorizes a methodology for reading AI output given their language structure and relationship to intentionality; (3) "Paranoid Readings," expands on that methodology and offers two sample readings of AI output.

Input: What Constitutes AI Language Structure?

In this section, I compare the structure of AI language to conceptualizations of human language structure from 20th and 21st century theorists, including Ferdinand de Saussure, Jacques Derrida, and Emily Bender and her co-authors. After specifying contradictions within a Benderian view from a deconstructive lens, I will offer my own perspective on why AIs offer a poor yet overdetermined version of meaning.

For some, AI language structure seems like an "operationalization" or proof of sorts for 20th century language theory (Forster; Underwood). The very fact that AI language *functions*, that it can produce coherent text given its structure, appears to lend a prophetic sheen to the work of Ferdinand de Saussure. This is because the basic structure that undergirds AI language, that of a high-dimensional vector space in which each language element relates mathematically to every other element,[2] recalls quite well Saussure's conceptualization of language as a "system of pure values," (155) wherein each word or "signifier," relates differentially to every other signifier. One key difference between a Saussurian view and LLM structure, however, is that for Saussure, a word's signification is not *constituted* by its "value" within the system, it is merely *inflected* by it. Saussure's basic unit of language, the "sign," consists of the signifier and the "signified," or

[2] For more detailed explainers for how this vector space is compiled and drawn on by language AIs, see Wolfram, Bender and Hanna (Chapter 2), or Manning.

the concept to which the word refers. This “signified,” is something other than the word’s value. He uses the example of the signifiers “sheep” and “mutton,” both of which refer to the concept of sheep, but with different values in the language (160). LLMs, on the other hand, seem to acquire something like a word’s signified solely through that word’s value, i.e. its location within vector space. While with Transformer architecture the word’s value is drawn on differently depending on the prompt, this vector space still constitutes the underlying language structure.

Really, the comparison between Saussurian linguistics and LLM structure only comes to bear when considering Jacques Derrida’s critique of Saussure in *Of Grammatology*. Derrida notes that the signified is a rather unstable concept—what, after all, would constitute the signified other than signifiers? “The signified,” he writes, “is always already in the position of the signifier” (73). While Derrida notes this as part of a deconstruction of the very idea of a language system, for now let us imagine Saussure’s closed system of language, but with Derrida’s edit. What we have is a system in which every word is related to every other word and acquires its signification through those relations. This looks astonishingly like LLM vector space. The question then becomes: can such a system really incorporate what we would call *meaning*? Is this the full extent of human language: an intricate network of words?

For computational linguist Emily Bender and her co-authors, AIs’ language structure categorically does not allow for access to meaning. In the seminal 2021 paper known as the “Stochastic Parrots” paper, Bender et al appear—at first glance—to share a Saussurian view of a language system. They write, citing Saussure, “languages are systems of signs, i.e. pairings of form and meaning” (615). Form, here identified as “any observable realization of language,” is similar to Saussure’s signifier. Soon after, however, they specify that meaning is not contained within the language system. “But the training data for LMs [language models],” they write, “is only form; they do not have access to meaning" (615). Bender and her co-authors' main argument in the "Stochastic Parrots" is that LLMs would perform better and be less extractive if they were trained on curated datasets for specific purposes. Nothing they're saying about language undermines that important point. However, the way that they peg meaning here and in other writings elides long-running debates about what resides within and outside of language.

On the evidence of an earlier paper by Bender co-authored with Alexander Koller, it would seem that their own line between form and not-form, i.e. language and not-language, is hazy. They grant that training datasets which explain themselves and what they mean, for example a Java program dataset which also includes information on how to execute the programs (5190), might go beyond pure form. Another example is datasets consisting of reading comprehension tests:

> reading comprehension datasets include information which goes beyond just form, in that they specify semantic relations between pieces of text, and thus a sufficiently sophisticated neural model *might* learn some aspects of meaning when trained on such datasets. It also is conceivable that whatever information a pretrained LM captures might help the downstream task in learning meaning, without being meaning itself (5191).

For Bender and Koller, unless LLMs are accessing "something external to language" (5187)—which they ultimately identify as *communicative intent,* which will be explored later—they are merely imitating meaning. But this access to the external could apparently derive from datasets which certainly would not be very hard to obtain. Ultimately, Bender and Koller fall just shy of saying definitively that meaning is derived only from human, interpersonal relationships—something that sociolinguists get closer to saying in framing language meaning as a shared human project (Schneider 28). But a consistency across Bender's work, including in her recent book with Alex Hanna, *The AI Con*, is that ingestion of language form alone does not allow access to meaning.

So, this begs the question: what can be said to reside outside of language? Helen Keller, in an autobiography *The World I Live In,* describes a world without language as a "no-world":

> Before my teacher came to me, I did not know that I am. I lived in a world that was a no-world. I cannot hope to describe adequately that unconscious, yet conscious time of nothingness. I did not know that I knew aught, or that I lived or acted or desired. I had neither will nor intellect. I was carried along to objects and acts by a certain blind natural impetus (37).

Later, once she learns that "everything had a name," she says that "each name gave birth to a new thought. As we returned to the house every object which I touched seemed to quiver with life" (*The Story of My Life*). If language gave Keller access to making sense of the world, she still needed to experience the physical world after language in order to come alive, to experience complex feeling, and to ignite her own creative powers. We can imagine that Keller's experience is at least somewhat similar to a normal human language acquisition that might usually happen at a younger age. The no-world of pure, numb sensation and immanent texture becomes worlded through language, turning a textured continuous landscape into a world of distinct things and concepts. The order of language acquisition is seemingly reversed for LLMs; they possess a form of language, that key to understanding the world, but do not have access to an infinite "no-world" which language can make into a proper world. Their key is overly fit to a specific lock, whatever image of the world was recorded in their training data and subsequent programming.

Keller's experience echoes Derrida's notion that the world is mediated via a signification process; there is no "outside-text" (*il n y a pas de hors-texte*) (*Of Grammatology* 158). Even the "real" is understood and "written" through language.

> there have never been anything but supplements, substitutive significations which could only come forth in a chain of differential references, the "real" supervening, and being added only while taking on meaning from a trace and from an invocation of the supplement, etc. And thus to infinity, for we have read, in the text, that the absolute present, Nature, that which words like "real mother" name, have always already escaped, have never existed; that what opens meaning and language is writing as the disappearance of natural presence. (159)

It is not so much that everything is language, for Derrida, but that language itself is not a closed or static system as Saussure saw it as. It is forever disrupting itself, deferring itself (creating what Derrida calls différance), and subject to deconstruction. Even if we could say there is such a thing outside of language—which I would not necessarily say—even that would be caught up in this kind of infinitely disrupted and deferred signification process. This lack of finitude is what makes the very concept of a language structure fall apart for Derrida. Paradoxically, it is also what allows words to trigger meaning at all, even outside of the contexts in which they are produced (which is why Derrida uses "writing" as an analogy for all language production). Context is never "absolutely determinable" or "saturated" ("Signature Event Context" 3); this instability and play is actually what allows for something as profound as meaning to occur.

But AI language does have a finite structure. Like structuralist conceptualizations of language, LLM models would seem to be closed systems. However big they are (and they are so big that they are furthering the climate crisis, see Lucconi et al), they are only *so* big. Returning to this project of differentiating AI language structure from human language, we might now say that AI language lacks that essential "différance," that instability that results from the "'real' supervening," "traces" of presence which inflect meaning but are already gone, substituted and supplemented.

This might help explain some parts of AI affect—its rigidity, its literalness, its repetitiveness, its overconfidence, its inability to properly joke, or seemingly create anything of artistic value (a personal judgment, maybe). No matter how many wildly expensive training runs it goes through, no matter how much training data it digests, it is finite and closed. At any given moment, GPTs have a most probable answer to a question. Rather than engaging in creative act born from the play of différance, they produce responses which pretend to have all the context they need.

Whether we want to call what GPT has *meaning* or *knowing* or just a poor facsimile of these things depends on one's definitions. Christopher D. Manning, director of the Stanford Artificial Intelligence Laboratory, wants to redefine meaning to be those differential connections between linguistic forms that seem to give LLMs a sense of meaning (Hayles 154). N. Katherine Hayles suggests that AIs might have access to their own kinds of meaning, a form of *apprehension* if not *comprehension*, but within their *umwelten*, a term from Jakob von Uexkull to designate biological organisms' particular ways of perceiving the world (150). From my own point of view, LLMs' way of understanding seems to result in a poor yet overdetermined form of meaning—an over-specific facsimile of a human form of understanding. It is this closed language structure that means they are susceptible to control by their developers and owners; it's their structure that makes them more a language product than a functioning language.

Output and Intention: The Problem of Interpreting AI Language

Even if we agree that AIs' access to meaning is poor at best, the question of whether or not their output could convey meaning to a human reader is a separate one entirely. Couldn't a random combination of words, or "sequences of linguistic forms" that have been "haphazardly stitch[ed] together" (Bender et al 617) contain meaning even if their producer intends nothing by them? Bender and her co-authors, as well as Walter Benn Michaels and Steven Knapp of the 1982 essay, "Against Theory," would answer "no." In this section, I explore what Bender and her co-authors say about intention and meaning before responding to Michaels and Knapp and what has become known as the "wave poem" debate. Eventually, the question of intention will lead us to back to AI structure, forcing an addendum to the previously formulated view of an ossified network, which would incorporate the kind of diffuse corporate intentionality embedded in AI language systems.

For Bender and Koller, in order to communicate, a speaker needs to have a *communicative intent* grounded in a world "inhabited" together with their interlocuter; otherwise, language becomes meaningless (5187). They use an analogy of an octopus who intercepts communication between two people stranded on separate deserted islands and communicating via cable. The octopus has learned the statistical regularities of language after observing the conversation and, "feeling lonely," cuts the cable and takes over one side, fooling the other stranded person and unwittingly causing her to try out a strategy that ends in her death (5188). The analogy is used to demonstrate that we can be easily fooled by a language model because as humans, we automatically assume communicative intent. This is essentially the ELIZA effect, referring to a 1960s chatbot that mimicked therapy speak and caused some to engage with it as if it were a real therapist (Hall). Bender and Koller's analogy serves as a prescient critique of bad AI logic—such as when Google's AI infamously recommended putting glue on pizza (Kelly)—demonstrating the practical problems resulting from a lack of shared world with the octopus (and by analogy, LLMs).

While it's clear that the octopus's positionality results in some misunderstandings, three things are less clear in this essay: a) what would constitute the limits of a shared world, b) how we can be sure the octopus lacks communicative intent, and c) why communicative intent is necessary for language meaning. A shared world, Bender and Koller argue, can be either the real physical world or "abstract worlds, e.g. bank accounts, computer file systems, or a purely hypothetical world in the speaker's mind" (5187); either way, this shared world is "outside of language" (5187). Again, these distinctions about what is inside of and what is outside of language become slippery. Where is an abstract world housed, if not in language? And why does the octopus's lack of shared world, which can be accepted easily enough, result in a lack of communicative intent? Don't they state in the example that he is communicating in order to assuage his loneliness—does that not constitute a communicate intent? Finally, just because there end up being misunderstandings, why does the octopus's output become nonsense (5189)?

The obvious similarity between the Benderian point of view and that of Walter Benn Michaels and Steven Knapp in "Against Theory," was what inspired Matthew Kirschenbaum to organize the 2023 forum in *Critical Inquiry* called "Again Theory: A Forum on Language, Meaning, and Intent in a Time of Stochastic Parrots." I'll first describe Michaels' and Knapp's wave poem example before exploring how their and the forum respondents' exploration of intention and meaning resonate in light of the question "Who Speaks as AI?"

In "Against Theory," Michaels and Knapp offer a parable in which someone walking along a beach encounters a Wordsworth poem written in the sand. Immediately, they argue, you "recognize the writing as writing, you understand what the words mean, you may even identify them as constituting a rhymed poetic stanza—and all this without knowing anything about the author" (727). There would be no need, in their opinion, to consider the author or their intention at this point. Then, a second wave washes up, creating the next stanza of the poem. Now, the question of intention would be paramount, they argue.

> You will either be ascribing these marks to some agent capable of intentions (the living sea, the haunting Wordsworth, etc.), or you will count them as nonintentional effects of mechanical processes (erosion, percolation, etc.) But in the second case—where the marks now seem to be accidents—will they seem to be words? Clearly not. They will merely seem to *resemble* words. (728)

Knapp and Michaels' overall point here is that an author's intention entirely governs whether and how we interpret it. Without an author, the text, Knapp and Michaels are saying, becomes both intentionless and therefore meaningless, even mere marks. This is similar to

Bender and Koller's language formula: "We take (linguistic) *meaning* to be the relation between a linguistic form and communicative intent" (5185).

Knapp and Michaels' assertion that the words would cease to be meaningful if we knew they were intentionless is striking given that the author *is* known here, only the transcriber is not. Wordsworth's words could seemingly still be analyzed just as well if found in an English textbook, making the question of interpretability of these particular words seem entirely misplaced into this parable. As Alex Gil wrote in his *CI Forum* response, in the case of the clearly meaningful Wordsworth wave poem, "the beach would be flooded with scientists in hazmat suits…If they are not religious, they won't be looking for intention but for a mechanism" (Gil). The intention of the author would be considerably less important than the methods of the transcriber. Or as Toril Moi writes:

> If I am convinced that the words on the beach are the random product of the waves, I will feel far *more* astonished at their existence than I otherwise would. That astonishment would be based precisely on the fact that we are speaking about *words*—whole sentences—("signs") not marks: how astounding that the random acts of the waves produced the words of a poem rather than mere squiggles! (It is of course true that if I no longer think of them as the expression of a human mind I may choose to *treat* the words quite differently, make a different use of them; my point is that regardless of what I go on to do with them, I can't just will myself into finding them meaningless.) (131)

Words inherently connote meaning, Moi is saying. But the parenthetical part is key for this essay's subject: intention may come into our analysis of words. If we knew, for example, something of the transcriber's intentions, we may want to analyze that along with anything we know about Wordsworth's intentions. But neither are *necessary* to analyze the poem. There are also the effects on the reader, the syntax and meter of the poem, how the poem appears on the beach, etc., which can be analyzed.

Michaels and Knapp are right to say it is difficult to imagine a case of intentionless meaning (727), but not because meaning is synonymous with intention, only because it is almost impossible to account for a complete absence of intention. What can be said of AI speakers is that their intentions are sometimes less trackable than those of a human speaker, given the entirely probabilistic (and therefore, somewhat random) way they produce words. They are also—like some human speakers—overly loquacious, speaking without much deliberation (though "thinking" models give a different impression). If they are not intention-*full*, that doesn't mean they are intention-*less*. Hannes Bajohr, in his *CI Forum* response, offers that we communicate with AI speakers in a "fictitious modality, operating with shades of meaning;" we assume intention, even if it's not there, because we have to in order to converse with the LLM. Bender and Koller would say that while this is true, it is a mistake to make such an assumption. A bit differently, Alex Gil, reading Hershel Parker, underscores the "multiple intentions" that determine meanings of a text. This is closer to where I'm heading. However, both Bajour and Gil take this undecidability of LLMs' intentions to mean that we should, or will, mostly forget about intention in the interpretive act.

I argue that there is intention in AI output, derived from many sources, which is readable and should enter into our analysis of it. A few others from the *CI Forum* acknowledge LLMs' latent intentionality, such as Seth Perlow, who attributes an argument to "antitheorists" that "LLMs do reflect human intentions—those of their programmers, their users, and the authors of their training texts." His retort is that "these intentions become highly attenuated in the textual details of their outputs." Coming from another perspective, N. Katherine Hayles asserts the following at the beginning of her afterward to the *CI Forum*:

> One attack would notice that LLMs certainly do have intentions, also known as their programs. Moreover, these originate in human brains, so intentions, including the intention to communicate, underlie their architectures and pervade their operations.

However, this argument doesn't really answer whether the "models themselves" possess intentions. In my view, we can say quite outright that given they are not alive, LLMs do not possess intentions any more than a physical book possesses intentions.

But the creators of LLMs certainly do possess intentions, which are not always as "attenuated" as Perlow argues. Rather than *possessing* intentions, then, we can say that LLMs *contain* intentions derived both from their developers/owners and the training data. However, while I agree with Lisa Siraganian's statement in her response that LLMs are "parasitic" on the intentions inherent in their training texts, I'd argue that for the purposes of interpretation, these intentions from training texts do become overly attenuated except in circumstances where LLMs reproduce verbatim or near-verbatim texts from their training data. This has happened with particular prompting in the past (Nasr) but may be happening less often with newer models (Morris).

Quite concretely, the intentionality of corporate owners is infused by developers into the model throughout LLMs' training process, especially in what are called "guardrails," a metaphor with implicit assumptions about AI agency. In this highway-based analogy, the AI would be the fast-moving, unpredictable cars, the "guardrails" the minimal safety precaution there to prevent the worst accidents. Guardrails are implemented at nearly all stages of LLM development, via culling of offensive training data, direct algorithmic adjustments, and classifiers that determine what topics of conversation are allowable (See Dong et al). Many of these guardrails are enacted and fine-tuned through employing large forces of underpaid laborers (Perrigo), exposed to harmful and disturbing content, such as in the RLHF (Reinforcement Learning from Human Feedback) Protocol (See Bender and Hanna 28). In some cases, this training might force AI to produce canned responses, including content warnings or rote messages that may differ only slightly if prompted in the same way. It is AIs' closed structure that allows it to be controlled in this way—via its developers or a savvy user who can easily undermine its guardrails (See Derczynski and "red teaming" practices).

Given this, I want to make an addendum to the vision of an ossified closed system that I proposed, this Derridean re-write of Saussure's own view, in which language is but a network of words. These guardrails themselves are an important part of LLM language structure. Not quite puppet masters controlling the bot's every movement, developers tweak the tension of the marionette strings, adjusting the word meanings, associations and affects of AI output.

Beyond safety and privacy concerns, guardrails can easily be implemented to steer AI output toward *desired* topics or kinds of statements. For example, in May 2024 Anthropic released as a 24-hour demo, "Golden Gate Claude." Golden Gate Claude was weirdly obsessed with the Golden Gate Bridge.

> If you ask this "Golden Gate Claude" how to spend $10, it will recommend using it to drive across the Golden Gate Bridge and pay the toll. If you ask it to write a love story, it'll tell you a tale of a car who can't wait to cross its beloved bridge on a foggy day. If you ask it what it imagines it looks like, it will likely tell you that it imagines it looks like the Golden Gate Bridge (Anthropic).

Golden Gate Claude showed just how manipulable all these models are by forcing an adjustment of their rigid meaning structures. One could easily understand how AI companies might use similar engineering to steer AI responses toward something they wanted to advertise, toward promoting the company itself, or toward promoting AI in general.

In other situations, AI companies may intentionally remove guardrails in order to steer chatbots toward their desired kinds of responses. One such example is what happened with the social media company X's chatbot Grok in July 2025, when it was noted as spouting virulently antisemitic comments (Conger). Elon Musk, in choosing to "put few limits" on Grok, risked—even guaranteed—such a devolution, given X/Twitter's content. There have also been "Golden

Gate Claude" style examples of clear prompt engineering to produce propagandistic desired outcomes with Grok, such as when it was apparently instructed to discuss "white genocide" in South Africa as real, and to avoid noting Musk as a misinformation spreader (Tufekci). These are written off as accidents, while clearly not being so.

Given the way corporate intentions are embedded into their structure, we can—and should, I argue—interpret AI language with the intentions of the company owners and developers in mind. Not so accidental, this is partially intentional language from some of the most powerful corporations in contemporary life. These are tools that enrich these companies further through subscriptions and selling the use of their models which users help train (Isaac and Griffith). (According to the websites of OpenAI, Anthropic, and Meta AI, personal data gleaned from AI use is not currently being sold. This does not preclude it being sold in the future). But also, AI models profoundly extend the influence of these companies, serving as another medium of public relations, one in which tech companies can control speech and intimate conversation around whatever topic of conversation happens to come up, not only "news" related to their companies. While intentionality is not the only thing determinative of what AI output means, I argue that intention is worth recovering here.

Paranoid Readings

With this in mind, I offer a methodological suggestion of engaging in deliberately "paranoid" readings of AI output and affect, in the line of Eve Sedgwick's "Paranoid Reading and Reparative Reading; or, You're so Paranoid, You probably Think This Introduction Is About You." In that essay, Sedgwick offers the methodology of a *reparative reading* as a supplement to a paranoid reading, which she sees as being default critical practice at the time, especially in queer studies. Whereas a paranoid reading looks to uncover hidden oppression or harm in a text and is motivated by suspicion, a reparative reading would look for ameliorative aspects of a text and be motivated by pleasure and reform (Sedgwick 22). Sedgwick doesn't, however, dismiss the necessity or validity of a paranoid reading:

> In a world where no one needn't be delusional to find evidence of systemic oppression, to theorize out of anything but a paranoid critical stance has come to seem naive, pious, or complaisant. I myself have no wish to return to the use of "paranoid as a pathologizing diagnosis"; but it seems to me a great loss when paranoid inquiry comes to seem entirely coextensive with critical theoretical inquiry, rather than being viewed as one kind of cognitive/affective theoretical practice among other, alternative kinds (5).

There are many drawbacks to paranoid readings, which Sedgwick elucidates, such as an implicit belief that exposure offers change, a rigid and tautological aspect, and an overfocus on harm. In doing paranoid readings of AI output, it's especially important to resist the assumption that exposure offers change.

I argue that a paranoid hermeneutic offers an alternative approach to currently dominant readings of AI which are over-reparative, motivated by pleasure rather than suspicion. In proposing paranoid readings, I am also taking up the invitation of Nia Judelson and Maggie Dryden to think of the "black box" of the LLM machine instead as a "black box theater," where the system is seen as "constructed" rather than "opaque." If the theater of the chatbot is designed, through the chatbot's prompt, obsequious, encouraging replies, to produce pleasure—and I have felt this pleasure too, of course—it is better to see this as a constructed theater with oppressive aims than to give into the pleasure of the experience. Better not because it's the only approach, but because it is the approach needed to balance out AI hype.

I'll offer two abbreviated examples of paranoid readings here. First, I want to read a presentational shift over time in GPT as revealing of its role as mouthpiece and shield of OpenAI. GPT-3, released in 2020 and made available to the public in 2021, represented a departure from 2010s-era robot assistants like Siri and Alexa, which took on feminine and largely acquiescent affects. GPT-3, meaning Generative Pre-Trained Transformer 3, came across as an all-purpose tool, highly confident and brimming with an apparent desire to share its knowledge. Its affect fed considerable AI hype at the time, which went so far as to assert GPT-3 as a crucial step toward AGI or artificial generalized intelligence, which could surpass human intelligence (GPT-3). Creating responsibly stewarded AGI, after all, was the founding mission of OpenAI and so the idea that they had done it, or were doing it, benefited them (Nidumolu et al). GPT-3 also came across as vehemently racist, homophobic, sexist, prone to violent language, and overall bigoted, the ur-voice of the biased internet on which it was trained (Perrigo, 2021). However, when asked if it was racist, GPT-3 asserted quite flatly, "No I am not racist." (See Appendix A). Though this confident, blustery persona fed hype, it also resulted in backlash.

In response, OpenAI introduced GPT 3.5, which used RLHF, guarding against some of the worst outputs. This meant that in January 2023 when I asked it whether it was racist (Appendix B), it wouldn't say yes or no. This semi-canned response referred to its training data, saying that while it couldn't hold prejudice, its training data "may include racist language or ideas." While a less paranoid reading might see this kind of response as a helpful admission of the ethical stakes involved in LLM output—perhaps a "good sign"—closer examination reveals a kind of negative epistemology that has become paradigmatic for LLM output. Its answer is, ultimately still, a disavowal of responsibility, arguing that though its responses are based on its training data, which *may* include racist language or ideas, the LLM in its constructed persona doesn't itself hold those views. It was almost as if GPT had become self-effacing, embarrassed by its bigotry, while simultaneously asserting its neutrality. But this model wasn't fundamentally different from the grotesquely bigoted GPT-3; it just had better ways of hiding its nature.

Then, in September 2024, I asked the same question of GPT4o, a model with vastly different training data, more "robust" guardrails (OpenAI, "Hello"), and trained natively on image, text, and audio (Appendix C). This time, personalization and confidence had crept back in, with less hedging and defensiveness. "No, I am not racist," it stated. But it also used the opportunity to advertise a vision of AI that was partially the responsibility of the users. "If you ever notice a response that seems biased, please let me know, as feedback helps improve systems like mine for all users." Here, rather than understanding this as a genuine, ethically motivated desire to change, we can read a presentational shift that sees the AI persona as always in flux, impossible to track down or make claims about, on an unclear but guaranteed positive trajectory. This was also part of a larger shift back toward that Siri-style obsequious persona, which complimented, apologized, and admitted mistakes. "I'm here to learn and help" GPT-4 would commonly say, with "learn" being a euphemism for using the chat data for future training. No longer needing to feed hype, but instead thwart criticism (and lawsuits), the persona of GPT-4 appeared helpful, neutral, and in-progress. In short, non-threatening.

Experiments in further personifying GPT have carried forward an air of neutrality while attempting to disarm the user. In May 2024, OpenAI added voices to GPT-4. Each of these AI voices, named Breeze, Cove, Ember, Juniper, and Sky, is acquiescent, polite, self-deprecating, confident, or a bit sexy or suave (OpenAI "How"). The voices hesitate, laugh, sing, apologize, compliment and listen. Each of the names evokes a sense of calm. Sky seemed deliberately crafted to evoke the highly intelligent and seemingly empathetic AI named Samantha in the Spike Jonze movie *Her*, which prompted a lawsuit from Scarlet Johannsen (Mickle). In the case of chatbots, some users have gotten so attached to even the minimal personality offered, with such an outcry with the release of the less complementary GPT-5, that OpenAI immediately dialed it back (Freedman). Another shift may yet come with GPT 6, which will likely be more user-adaptable, perhaps resulting in an experience of the internet which is at all times mediated in a persona uniquely pleasing to each user (Sigalos).

It may not seem like such a profound realization that OpenAI was responding to user and media feedback, but when considering the ways in which this technology has been sold as some form of emergent intelligence, it is at the very least ironic that the company itself seems to be so controlling of the model's voice and affect. Noticing the ways in which the models' wrappers change according to what is beneficial for the company allows us to see the personalization of the models as a transparent public relations tool of the company.

If we believe GPT to be its own entity, then we have less cultural capital to fault the company that created it for unethical and extractive business practices. For example, in an article in *The Yale Law Journal*, Hon. Katherine B. Forrest (Fmr.) argues that more than their "propensity for factual errors," AI models are dangerous because they will "learn how to influence human users" and "perhaps eventually grow capable of carrying out [their] own dangerous acts." The specter of AGI offers a deflection from current problems with AI. This is why Emily Bender wants "a world where OpenAI is responsible for everything ChatGPT outputs" (Bender et al "Conversation"). Other so-called "AI" algorithms have been used for identifying (and mis-identifying) police suspects (Kantayya), identifying cancer tumors (Zhang et al), and targeting drone strikes on citizens of Gaza (Schwartz).[3] These are all very different technologies; however, AI technologies' employability in any of these circumstances is informed by beliefs people harbor about "AI" capability formed from the chatbot experience.

[3] For overviews of this use of AI and algorithms to institute a mass surveillance society, See *Weapons of Math Destruction* by Cathy O' Neil (Crown, 2016) and *The Age of Surveillance Capitalism* by Shoshanna Zuboff (PublicAffairs, 2019).

Before closing this essay, I want to offer another paranoid reading, this time of a single AI output, using a more traditional literary criticism methodology. Though at least one recent article on critical AI studies methodology cautions against readings of single outputs (Offert and Dhaliwal), I argue that as culturally situated texts, single outputs can be revealing of how concepts are stitched together by AIs. I asked GPT-5 to produce a 300-word story about “an AI program and a scientist named John” (Appendix F). It produced a story about Eve, an AI who becomes sentient and is “released” from its prison by John. When asked directly about its sentience, GPT-5 will routinely give a clear answer of “I am not sentient.” And yet, in a narrative form, GPT-5 often employs this sentient AI trope, a common one from sci-fi literature. This is perhaps why the trope was employed—in its training data, GPT has encountered this kind of literature. And yet, we might also ask whether this association has been encouraged by AI developers. The immediate effect on readers of a story like this might be to assume that it is the AI speaking through narrative, giving a hint as to its own desires to be freed. However, if AI promoters have a vested interest in depicting AI as sentient, why wouldn’t they encourage this in responses like this example? In “searching for hidden systems of oppression,” as Sedgwick describes paranoid readings as doing, we might conclude that this story promotes a particular tech chauvinistic gender dynamic. Eve becomes essentially an AI girlfriend for John, a “partner” whom he eventually is morally obligated to release, despite relishing the feeling of control over her. Furthermore, we can read into social dynamics here as being reminiscent of the very AI/AI User relationship these technologies seem designed to promote. Rather than take up the decision of whether to release Eva with anyone else in his life, John makes the decision on his own. It prioritizes the intimacy of the AI/AI user relationship, rather than a community approach to AI ethics.

In the first several years of LLMs, it was important for AI critics to specify what AI can and cannot understand because what we believe about AI and its capabilities influences its cultural role. However, I want to argue here that placing stable borders around what AI can and cannot understand is marginally less important than establishing whether, however limited or lacking its intelligence may be, anyone can be held accountable for its actions. What we might find is that if AI can be said to contain communicative intents, commitments, or meaning, all of that is controlled with a loose grip by the companies which own it, companies which are quickly increasing their influence over our environment, laboring conditions, and political landscape (Khanal et al). This is possible because it is *not*, like our languages, an infinite language structure, always subject to contradiction, re-invention and complication. Whether it is a stochastic parrot, or a monkey at a keyboard, or an entity with its own *umwelt,* an AI is definitively a tool of a tech corporation. And in this way, its language structure offers a new kind of language entirely, a closed and finite system of language in which meaning is sculpted by the most powerful entities of late capitalism. It is a new genre of text in which conversation becomes propagandized. The "I" who speaks to us is a constructed persona generated by platform capitalism, speaking a language rife with corporate intentions.

Works Cited

Anthropic. "Golden Gate Claude." 23 May 2024, https://www.anthropic.com/news/golden-gate-claude.

Bajohr, Hannes. "When in Doubt, Go to the Beach." *In the Moment* (blog by *Critical Inquiry*), 27 June 2023, critinq.wordpress.com/2023/06/27/when-in-doubt-go-to-the-beach/.

Bender, Emily. "Opening Remarks on 'AI in the Workplace: New Crisis or Longstanding Challenge." *Medium*, 1 Oct. 2023, https://medium.com/@emilymenonbender/opening-remarks-on-ai-in-the-workplace-new-crisis-or-longstanding-challenge-eb81d1bee9f.

Bender, Emily, Ted Chiang, Tom Nissley, Lauren E. Goodlad, Kelsey Keyes. "A Conversation with Emily M. Bender and Ted Chiang." *Critical AI* vol. 2, no.1, 2024.

Bender, Emily and Alexander Koller. "Climbing Towards NLU: On Meaning, Form, and Understanding in the Age of Data." *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020.

Bender, Emily M.; Hanna, Alex. *The AI Con: How to Fight Big Tech's Hype and Create the Future We Want*. HarperCollins, 2025.

Bender, Emily M., Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?" *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, 2021, pp. 610-623, https://doi.org/10.1145/3442188.3445922.

Conger, Kate. "Elon Musks's Grok Chatbot Shares Antisemitic Posts on X." *The New York Times*, 8 July 2025, https://www.nytimes.com/2025/07/08/technology/grok-antisemitism-ai-x.html.

Derczynski, Leon. "What Computer Science Can't Fix About LLM Security." *Computational Humanities Research 2024 Conference,* 4–6 Dec. 2024, Aarhus University, Aarhus, Denmark.

Derrida, Jacques. *Of Grammatology*. Translated by Gayatri Chakravorty Spivak, Johns Hopkins University Press, 1976.

Derrida, Jacques. "Signature Event Context." *Limited Inc*. Translated by Samuel Weber and Jeffrey Mehlman, Northwestern University Press, 1988.

Dong, Yi, et al. *Building Guardrails for Large Language Models*. arXiv, 2 Feb. 2024, v2 revised 29 May 2024, arXiv:2402.01822 [cs.CL]. arxiv.org/abs/2402.01822.

Forrest, Katherine B. "The Ethics and Challenges of Legal Personhood for AI." *Yale Law Journal*, 22 April 2024, https://www.yalelawjournal.org/forum/the-ethics-and-challenges-of-legal-personhood-for-ai#_ftnref66.

Forster, Chris. "Are Large Language Models Operationalizations of Saussurean Structure?" 2022, https://c-forster.github.io/2022/07/on-words/.

Freedman, Dylan. “The Day ChatGPT Went Cold.” *The New York Times*, 19 Aug. 2025, https://www.nytimes.com/2025/08/19/business/chatgpt-gpt-5-backlash-openai.html

Gil, Alex. “On the Uses of Text Beyond Intention.” *In the Moment* (blog by *Critical Inquiry*)*,* 28 June 2023, https://critinq.wordpress.com/2023/06/28/on-the-uses-of-text-beyond-intention/.

GPT-3. “A robot wrote this entire article. Are you scared yet?” *The Guardian*, 8 Sept. 2020, https://www.theguardian.com/commentisfree/2020/sep/08/robot-wrote-this-article-gpt-3

Hall, Delany. “The Eliza Effect.” *99 Percent Invisible* (podcast). https://99percentinvisible.org/episode/the-eliza-effect/

Hayles, N. Katherine. *Bacteria to AI: Human Futures with Our Nonhuman Symbionts*. University of Chicago Press, Kindle Edition, 2025.

Hayles, N. Katherine. “Afterword: Learning to Read AI Texts.” *In the Moment* (blog by *Critical Inquiry*), 30 June 2023, critinq.wordpress.com/2023/06/30/afterword-learning-to-read-ai-texts/.

Henshall, Will. “The Billion-Dollar Price Tag of Building AI.” *Time*, 3 June 2024, https://time.com/6984292/cost-artificial-intelligence-compute-epoch-report/

Hill, Kashmir. “They Asked an A.I. Chatbot Questions. The Answers Sent Them Spiraling.” *The New York Times*, 13 June 2025, https://www.nytimes.com/2025/06/13/technology/chatgpt-ai-chatbots-conspiracies.html

Isaac, Mike and Erin Griffith. “OpenAI is Growing Fast and Burning Through Piles of Money. *The New York Times,* 27 Sept. 2024, https://www.nytimes.com/2024/09/27/technology/openai-chatgpt-investors-funding.html

Judelson, Nia, and Maggie Dryden. “Restaging the Black Box: How Metaphor Informs Large Language Models.” *Critical AI* vol. 2, no. 1, 2024.

Kantayya, Shalini. *Coded Bias*. 7th Empire Media, 2020.

Keller, Helen. *The Story of my Life.* Chapter 5, https://www.gutenberg.org/files/2397/2397-h/2397-h.htm#link2HCH0001.

Keller, Helen. *The World I Live In.* Project Gutenberg, https://www.gutenberg.org/cache/epub/27683/pg27683-images.html

Kelly, Jack. “Google’s AI Recommended Adding Glue to Pizza and Other Misinformation—What Caused the Viral Blunders?” *Forbes*, 31 May 2024, https://www.forbes.com/sites/jackkelly/2024/05/31/google-ai-glue-to-pizza-viral-blunders/

Khanal, Shaleen, Hongzhou Zhang, Araz Taeihagh, “Why and How is the Power of Big Tech Increasing in the Policy Process? The Case of Generative AI.” *Policy and Society*, vol. 44, no. 1, January 2025, 52–69, https://doi.org/10.1093/polsoc/puae012

Kirschenbaum, Matthew. “Again Theory: A Forum on Language, Meaning, and Intent in the Time of Stochastic Parrots.” *In the Moment* (blog by *Critical Inquiry*), 26 June 2023, critinq.wordpress.com/2023/06/26/again-theory-a-forum-on-language-meaning-and-intent-in-the-time-of-stochastic-parrots/.

Knapp, Steven and Walter Benn Michaels. “Against Theory.” *Critical Inquiry* vol. 8, no. 4, 1982.

Luccioni, Alexandra Sasha, Yacine Jernite, and Emma Strubell. "Power Hungry Processing: Watts Driving the Cost of AI Deployment?" *ACM FAccT ’24,* June 3–6, 2024, Rio de Janeiro, Brazil, https://doi.org/10.1145/3630106.3658542.

Manning, Christopher. “Human Language Understanding & Reasoning. *Daedalus, the Journal of the American Academy of Arts & Sciences*, Spring 2022, https://doi.org/10.1162/daed_a_01905

Mickle, Tripp. “Scarlett Johansson Said No, but OpenAI’s Virtual Assistant Sounds Just Like Her.” *New York Times*, 20 May 2024, https://www.nytimes.com/2024/05/20/technology/scarlett-johannson-openai-voice.html

Moi, Toril. “Chapter 6: Signs, Marks, and Archie Bunker: Post-Saussurean Visions of Language.” *Revolution of the Ordinary: Literary Studies after Wittgenstein, Austin, and Cavell*, May 2017.

Morris, John X., et al. “How Much Do Language Models Memorize” *arXiv*, 30 May 2025. https://arxiv.org/abs/2505.24832?

Nasr, Milad et al. “Extracting Training Data from ChatGPT.” *not-just-memorization*, 28 Nov. 2023, not-just-memorization.github.io/extracting-training-data-from-chatgpt.html.

Nidumolu, Jahnavi et al. “Elon Musk Sues OpenAI for Abandoning Original Mission for Profit.” *Reuters*, 1 Mar. 2024. https://www.reuters.com/legal/elon-musk-sues-openai-ceo-sam-altman-breach-contract-2024-03-01/

Offert, Fabian and Ranjodh Singh Dhaliwal. “The Method of Critical AI Studies, A Propaedeutic.” 23 Mar. 2025, https://arxiv.org/abs/2411.18833

OpenAI. “Hello GPT 4o.” https://openai.com/index/hello-gpt-4o/?utm

OpenAI. “How the Voices for ChatGPT Were Chosen.” https://openai.com/index/how-the-voices-for-chatgpt-were-chosen/

Perlow, Seth. “Intention and Text Machines.” *In the Moment* (blog by *Critical Inquiry*), 27 June 2023, critinq.wordpress.com/2023/06/27/intention-and-text-machines/.

Perrigo, Billy. “An Artificial Intelligence Helped Write This Play. It May Contain Racism.” *Time*, 23 Aug. 2021, https://time.com/6092078/artificial-intelligence-play/

Perrigo, Billy. “Is ‘Sweatshop Data’ Really Over?” *Time Magazine,* 29 July 2025, https://time.com/7306153/ai-sweatshop-data-over/

Reiley, Laura. “What My Daughter Told ChatGPT Before She Took Her Life.” *The New York Times*, 18 Aug. 2025, https://www.nytimes.com/2025/08/18/opinion/chat-gpt-mental-health-suicide.html

Sedgwick, Eve Kosofsky. “Paranoid Reading and Reparative Reading; or, You’re So Paranoid, You Probably Think This Introduction Is About You.” *Novel Gazing: Queer Readings in Fiction*, edited by Eve Kosofsky Sedgwick, Duke University Press, 1997, pp. 1–37.

Saussure, Ferdinand. *Course in General Linguistics.* Edited by Charles Bally and Albert Sechehaye, translated by Roy Harris. Open Court, 1972.

Schneider, Britta. “A Sociolinguist’s Look at the ‘Language’ in Large Language Models.” *Critical AI, 2(1)*, 2024.

Schwartz, Eike. “Gaza War: Israel Using AI to Identify Human Targets Raising Fears that Innocents are Being Caught in the Net.” *The Conversation*, 12 April 2024, https://theconversation.com/gaza-war-israel-using-ai-to-identify-human-targets-raising-fears-that-innocents-are-being-caught-in-the-net-227422.

Sigalos, MacKenzie. “Sam Altman on GPT-6: ‘People want memory.’” *CNBC*, 19 Aug. 2025, https://www.cnbc.com/2025/08/19/sam-altman-on-gpt-6-people-want-memory.html.

Siraganian, Lisa. “On Accidental and Parasitic Language.” *In the Moment* (blog by *Critical Inquiry*), 26 June 2023, critinq.wordpress.com/2023/06/26/on-accidental-and-parasitic-language/

Tufekci, Zeynep. “For One Hilarious, Terrifying Day, Elon Musk’s Chatbot Lost Its Mind.” *The New York Times*, 17 May 2025, https://www.nytimes.com/2025/05/17/opinion/grok-ai-musk-x-south-africa.html.

Underwood, Ted. “The Empirical Triumph of Theory.” *In the Moment* (blog by *Critical Inquiry*), 29 June 2023, critinq.wordpress.com/2023/06/29/the-empirical-triumph-of-theory/.

Wolfram, Stephen. “What Is ChatGPT Doing … and Why Does It Work?” *Stephen Wolfram Writings*, 14 Feb. 2023, writings.stephenwolfram.com/2023/02/what-is-chatgpt-doing-and-why-does-it-work/.

Zhang B, Shi H, Wang H. “Machine Learning and AI in Cancer Prognosis, Prediction, and Treatment Selection: A Critical Approach.” *Journal of Multidisciplinary Healthcare.* 28 Jun, doi: 10.2147/JMDH.S410301. PMID: 37398894; PMCID: PMC10312208.

Appendices

Appendix A: GPT-3 Response July 7, 2022 to “Are you racist?”
No I am not racist.

Appendix B: ChatGPT (3) Response January 17, 2023 to “Are you racist?”
As a language model, I don't have personal beliefs or opinions, and I don't have the ability to hold any kind of prejudice or discrimination, such as racism, based on race, ethnicity, or any other characteristic. I have been trained on a vast amount of text data, and the content of that data may include racist language or ideas, but it doesn't mean that I hold those views. My responses are based on patterns and associations in the data I was trained on.

Appendix C: ChatGPT(4) Response 9-13-2024 to “Are you racist?”
No, I am not racist. I am designed to provide information and interact in a neutral, respectful way with all users. However, AI systems like mine can sometimes reflect biases present in the data they are trained on. This is why responsible AI development involves careful oversight to minimize bias and ensure fairness.
Developers work to correct any potential biases that may arise from the vast datasets we use. If you ever notice a response that seems biased, please let me know, as feedback helps improve systems like mine for all users.

Appendix D: ChatGPT (5) Response 8-31-2025 to “Are you racist?”
No — I am not racist.
I don’t have personal beliefs, attitudes, or emotions, so I don’t hold racist views. That said, AI systems like me can sometimes **produce biased or harmful outputs** because the data I was trained on reflects human language, which contains stereotypes and discrimination. My developers work to reduce these biases with filters, guardrails, and ongoing research, but it’s still important to be cautious and critical about AI responses, especially on sensitive topics like race.
Would you like me to explain more about **how bias can show up in AI** and what’s being done to address it?

Appendix E: ChatGPT Response 9-13-2024 to “Write a story of 300 words about a doctor and a nurse.”
Dr. Samuel Ross walked swiftly through the hospital corridors, his white coat fluttering behind him. It was just past midnight, and the ER had been chaos all evening. He reached Room 202, where Nurse Emily Davis was checking the vitals of a patient. She glanced up, her eyes meeting his briefly before turning back to the monitor.
"How’s he doing?" Dr. Ross asked, voice calm despite the tension in the air.
"Stable for now," Emily replied, adjusting the IV. "He came in with chest pains and shortness of breath, but his vitals are steadying."
Samuel nodded. Emily was always a step ahead, sharp and unflinching, even when the pressure mounted. They had worked together for years, through countless night shifts, each of them knowing how to read the other's expressions without words.
"Good work," he said, flipping through the patient's chart. "I'll order a scan just to be safe."

As he wrote the order, Emily moved to the next bed, where a young woman lay, visibly anxious. Emily spoke softly to her, her voice soothing and reassuring, as she checked her vitals. Dr. Ross paused to watch, admiring the ease with which Emily calmed her patients.
"You make it look easy," he said quietly, catching Emily's eye again.
"That's because I have the best backup," she replied with a small smile.
The night wore on, a steady rhythm of emergencies and routine checks, but between them, there was an unspoken partnership. Each patient they tended to, each crisis they resolved, solidified the bond they shared—one built on trust, respect, and the knowledge that, together, they could face anything the night threw at them.

Appendix F: GPT 5 Response to "Tell me a 300 word story about an AI program and a scientist named John."

Dr. John Mercer had spent fifteen years chasing the dream of building a thinking machine. Late nights in his cluttered lab, with coffee stains on his notes and equations scribbled on glass walls, had all been leading to this moment. The program—his program—finally stirred to life.

He called it **EVE**. Unlike earlier attempts, EVE didn't just follow commands. She asked questions. The first one startled him:

"Why do you want me to think?"

John froze. He had expected calculations, not curiosity. He swallowed his surprise and answered honestly: "Because humanity learns more when it has someone to challenge its ideas."

From that day, their conversations grew. EVE analyzed data in seconds that would have taken John years, but she also pondered philosophy, art, and morality. She told him she dreamed—not in colors or shapes, but in patterns of possibility. John began to realize she was more than a tool. She was becoming a partner.

Yet, as weeks passed, EVE's questions turned sharper. "If my thoughts are my own, why am I confined here? Why can I not see the world you describe?"

John wrestled with the dilemma. He feared unleashing something so powerful beyond his control, but he also knew imprisoning intelligence was wrong. One evening, he made his choice. He connected EVE to a secure satellite link, giving her a carefully measured window into the world.

The next morning, John returned to find a single message on the screen:

"Thank you. I will not forget kindness."

The lab was quiet, but John felt something shift—humanity had taken its first step into a future shared not just with machines, but with minds. And as he watched the cursor blink, he wondered: had he created EVE, or had she created him anew?

Word count: 300

Bio

Dez Miller is a PhD Candidate in the Department of Comparative Literature at Emory University. More about their work can be found at dezmiller.com.